

# LitinkAI Platform - Complete Software Architecture

---

## Executive Summary

LitinkAI is an AI-powered interactive platform that transforms books, articles, and documentation into engaging multimedia content for **learning**, **content creation**, and **entertainment**. The platform uses a microservices architecture with FastAPI backend, React frontend, and multiple AI services for content generation.

---

## Table of Contents

1. [System Overview](#)
  2. [Architecture Principles](#)
  3. [Technology Stack](#)
  4. [System Components](#)
  5. [Data Flow](#)
  6. [Security Architecture](#)
  7. [Deployment Architecture](#)
  8. [Scaling Strategy](#)
  9. [Documentation Index](#)
- 

## System Overview

### Vision

Transform static text content into interactive, multimedia learning and entertainment experiences using AI.

### Core Capabilities

- **Content Ingestion:** Upload books (PDF, DOCX, EPUB), articles, or AI prompts
- **AI Processing:** Generate scripts, scenes, characters, and plot overviews
- **Multimedia Generation:** Create images, audio, video, and interactive elements
- **Three Modes:**
  - **Learning Mode:** Educational content with quizzes and interactive lessons
  - **Creator Mode:** Professional content creation tools with plot management
  - **Entertainment Mode:** Interactive storytelling with character-driven narratives

### Key Differentiators

- Book-first approach (vs. course-first like Coursera/Udemy)
  - Multi-modal AI content generation
  - Tiered subscription model with 40-80% profit margins
  - Blockchain-based achievement credentials (NFT badges)
  - Microservices architecture for scalability
-

# Architecture Principles

## 1. Separation of Concerns

- Frontend handles presentation and user interaction
- Backend manages business logic and orchestration
- AI services handle content generation
- Database manages persistent state
- Message queue handles async processing

## 2. Scalability

- Horizontal scaling through microservices
- Async task processing with Celery
- Caching layer with Redis
- CDN for static assets
- Database read replicas for high traffic

## 3. Reliability

- Automatic fallback for AI models
- Circuit breaker pattern for external services
- Retry logic with exponential backoff
- Health checks and monitoring
- Graceful degradation

## 4. Security

- JWT-based authentication
- Row-level security in Supabase
- API rate limiting per tier
- Input sanitization
- Secure file upload validation

## 5. Cost Optimization

- Tier-based model selection
- Real-time cost tracking
- Usage monitoring and alerts
- Automatic model fallback to cheaper alternatives
- Caching to reduce API calls

---

# Technology Stack

## Frontend

```
Framework: React 18.3.1 + TypeScript 5.5.3  
Build Tool: Vite 5.4.2
```

Styling: TailwindCSS 3.4.1  
State Management: Zustand 4.5.0  
Routing: React Router DOM 6.22.0  
HTTP Client: Axios 1.6.7  
UI Components: Lucide React 0.344.0  
Video Player: Video.js (for interactive playback)  
Testing: Vitest 3.2.4 + Testing Library

## Backend

Framework: FastAPI 0.115.9  
Server: Uvicorn 0.34.3  
Language: Python 3.11+  
API Documentation: OpenAPI (Swagger/ReDoc)  
Validation: Pydantic 2.11.7  
Database ORM: SQLAlchemy 2.0.41 (minimal usage)  
Task Queue: Celery 5.3.4  
Message Broker: Redis 5.0.1  
Worker Monitor: Flower 2.0.1

## Database & Storage

Primary Database: Supabase PostgreSQL 15  
Authentication: Supabase Auth  
File Storage: Supabase Storage  
Vector Database: pgvector (for RAG)  
Cache: Redis 7  
Session Store: Redis

## AI & ML Services

LLM Router: OpenRouter (100+ models)  
Primary LLMs:  
- Free: Llama 3.2 3B, DeepSeek Chat  
- Premium: GPT-4o, Claude 3.5 Sonnet  
Image Generation: ModelsLab v7 (Imagen 4, SeeDream)  
Video Generation: ModelsLab v7 (VE02, VE03)  
Voice Synthesis: ElevenLabs v2  
Video Editing: FFmpeg (server-side)  
Speech-to-Text: OpenAI Whisper (future)

## Infrastructure


Containerization: Docker + Docker Compose  
Reverse Proxy: Traefik (suggested)  
Email Service: Mailgun (production) / Mailpit (dev)  
Payment Processing: Stripe  
Blockchain: Algorand SDK (NFT minting)  
Monitoring: Built-in health checks  
Deployment: Render (current)


## Development Tools

Testing: Pytest 7.4.3, Vitest  
Code Quality: Black, isort, Flake8, ESLint  
API Testing: HTTPx 0.27.0  
Debugging: debugpy (remote debugging)  
Version Control: Git  
CI/CD: (to be configured)

## System Components

### High-Level Architecture





## Component Descriptions

### 1. Frontend (React + TypeScript)

- **Location:** `/src`
- **Entry Point:** `src/main.tsx`
- **Key Features:**
  - Three distinct modes (Learning, Creator, Entertainment)
  - Real-time video generation status tracking
  - Interactive video player with AI chat overlay
  - Subscription management and usage tracking
  - Plot overview and character management
  - Multi-script support with versioning

### Key Components:

```

src/
└── components/
  
```

```

    ├── Admin/          # Admin dashboard components
    ├── Audio/         # Audio generation and timeline
    ├── Images/        # Image generation panel
    ├── Plot/          # Plot overview and characters
    ├── Script/        # Script generation panel
    ├── Subscription/ # Subscription tiers and usage
    ├── Video/         # Video production components
    └── VideoGeneration/ # Pipeline visualization
    contexts/          # React contexts (Auth, Theme, etc.)
    hooks/             # Custom React hooks
    pages/             # Page components (Dashboard, Profile, etc.)
    services/          # API client services
    types/             # TypeScript type definitions

```

## 2. API Gateway (FastAPI)

- **Location:** `/backend/app`
- **Entry Point:** `backend/app/main.py`
- **Port:** 8000
- **Features:**
  - RESTful API with OpenAPI documentation
  - JWT authentication via Supabase
  - CORS middleware for frontend
  - Rate limiting per subscription tier
  - Health check endpoint

### API Routes:

```

/api/v1/
    ├── auth/          # Authentication (register, login, verify)
    ├── books/         # Book management and upload
    ├── chapters/      # Chapter operations and AI content
    ├── ai/            # AI generation endpoints
    ├── plots/         # Plot overview management
    ├── characters/   # Character profiling
    ├── subscriptions/ # Subscription and billing
    ├── admin/         # Admin analytics and monitoring
    ├── merge/         # Video merging operations
    ├── payments/      # Stripe payment processing
    └── nfts/          # Blockchain badge minting

```

## 3. Core Services (Python)

### OpenRouterService

- Intelligent LLM routing based on subscription tier
- Automatic fallback to cheaper models
- Cost tracking and optimization

- Support for 100+ AI models

### AIService

- Multi-provider support (OpenAI, DeepSeek)
- Content generation (scripts, summaries, quizzes)
- Token management and sanitization
- Chapter analysis and metadata generation

### RAGService (referenced in docs)

- Vector embeddings with pgvector
- Context-aware content retrieval
- Chapter relationship mapping
- Enhanced plot generation context

### PlotService (described in architecture docs)

- Plot overview generation
- Character profiling with archetypes
- Story structure analysis
- Integration with OpenRouter

### SubscriptionManager (referenced in code)

- Tier-based access control
- Usage tracking and limits
- Cost monitoring and alerts
- Stripe integration

### ModelFallbackManager (referenced in code)

- Circuit breaker pattern
- Automatic model switching
- Retry logic with backoff
- Health monitoring

## 4. Async Task Processing (Celery)

- **Workers:** Defined in `backend/app/tasks/`
- **Broker:** Redis
- **Monitor:** Flower (port 5555)

### Task Modules:

```
app/tasks/
├── celery_app.py      # Celery configuration
└── image_tasks.py     # Character and scene image generation
└── video_tasks.py     # Video generation and processing
└── audio_tasks.py     # Audio/voice synthesis
└── merge_tasks.py     # Video merging with FFmpeg
```

```
└── lipsync_tasks.py    # Lip-sync processing  
└── blockchain_tasks.py # NFT minting on Algorand
```

**Task Flow:**

1. API endpoint queues task → Redis
2. Celery worker picks up task
3. Worker calls AI service or processes media
4. Worker updates database with progress
5. Worker notifies frontend via polling/webhooks

**5. Database (Supabase PostgreSQL)****Core Tables:**

```
-- Authentication & Users  
auth.users  
profiles  
  
-- Content Management  
books  
chapters  
scripts  
scene_descriptions  
  
-- AI Generations  
image_generations  
audio_generations  
video_generations  
merge_jobs  
  
-- Plot & Characters  
plot_overviews  
characters  
character_archetypes  
  
-- Subscriptions & Usage  
user_subscriptions  
usage_logs  
cost_aggregations  
  
-- Gamification  
badges  
user_badges  
nfts
```

**Key Features:**

- Row-Level Security (RLS) policies

- Vector embeddings (pgvector)
- Real-time subscriptions
- Automatic timestamps
- Foreign key constraints

## 6. AI Service Integrations

### **OpenRouter** (LLM Router)

- Base URL: <https://openrouter.ai/api/v1>
- Models: 100+ (GPT, Claude, Llama, DeepSeek, etc.)
- Cost Tracking: Per-request token usage
- Fallback: Automatic model switching

### **ModelsLab v7** (Image/Video)

- Image: Imagen 4 Ultra/Fast, SeeDream-4
- Video: VEO2, VEO3-audio, SeeDance
- Async Generation: Webhook callbacks
- Cost: ~\$0.002-\$0.10 per asset

### **ElevenLabs** (Voice Synthesis)

- Model: eleven\_multilingual\_v2
- Voice Cloning: Professional tier
- Streaming: Real-time audio generation
- Cost: ~\$0.30/1K characters

### **FFmpeg** (Video Processing)

- Merging: Concatenate video segments
- Transcoding: Format conversion
- Effects: Transitions, overlays
- Server-side processing

## Data Flow

### 1. Book Upload & Processing Flow

```

sequenceDiagram
    participant User
    participant Frontend
    participant API
    participant FileService
    participant AIService
    participant Database

    User->>Frontend: Upload book file
    Frontend->>API: POST /api/v1/books/upload
  
```

```

API-->FileService: Extract text content
FileService-->API: Raw text
API-->AIService: Generate chapters
AIService-->API: Chapter structure
API-->Database: Save book & chapters
Database-->API: Confirmation
API-->Frontend: Book preview
Frontend-->User: Display chapters

```

## 2. Script Generation Flow

```

sequenceDiagram
    participant User
    participant Frontend
    participant API
    participant OpenRouter
    participant RAGService
    participant Celery
    participant Database

    User->>Frontend: Request script generation
    Frontend-->API: POST /api/v1/ai/generate-script-and-scenes
    API-->Database: Check subscription limits
    Database-->API: Usage check passed
    API-->RAGService: Get chapter context
    RAGService-->API: Enhanced context
    API-->OpenRouter: Generate script (tier-based model)
    OpenRouter-->API: Generated script
    API-->Database: Save script
    API-->Celery: Queue scene image tasks
    API-->Frontend: Script response
    Celery-->Frontend: Async progress updates

```

## 3. Video Generation Pipeline

```

graph TD
    A[User Requests Video] --> B[Create video_generation record]
    B --> C[Queue Script Generation]
    C --> D[Generate Script with OpenRouter]
    D --> E[Parse Scenes & Characters]
    E --> F[Queue Character Image Tasks]
    F --> G[Generate Character Images - ModelsLab]
    G --> H[Queue Scene Image Tasks]
    H --> I[Generate Scene Images - ModelsLab]
    I --> J[Queue Audio Tasks]
    J --> K[Generate Audio - ElevenLabs]
    K --> L[Queue Video Tasks]
    L --> M[Generate Scene Videos - ModelsLab]
    M --> N[Queue Lip-Sync Tasks]

```

```
N --> O[Sync Audio & Video]
O --> P[Queue Merge Task]
P --> Q[Merge All Videos – FFmpeg]
Q --> R[Upload Final Video]
R --> S[Update Status: Complete]
S --> T[Notify User]
```

## 4. Subscription & Cost Tracking Flow

```
sequenceDiagram
    participant User
    participant Frontend
    participant API
    participant Stripe
    participant SubscriptionManager
    participant CostTracker
    participant Database

    User->>Frontend: Select subscription tier
    Frontend->>API: POST /api/v1/subscriptions/checkout
    API->>Stripe: Create checkout session
    Stripe->>API: Session URL
    API->>Frontend: Redirect to Stripe
    User->>Stripe: Complete payment
    Stripe->>API: Webhook: subscription.created
    API->>SubscriptionManager: Update user tier
    SubscriptionManager->>Database: Save subscription
    API->>Frontend: Webhook confirmation

    Note over User,Database: During AI usage
    User->>API: Generate content
    API->>CostTracker: Track API cost
    CostTracker->>Database: Log usage
    Database->>CostTracker: Check limits
    CostTracker->>API: Usage within limits
```

---

## Security Architecture

### Authentication & Authorization

#### Supabase Auth:

- JWT-based authentication
- Email verification required
- Password reset flow
- Refresh token rotation
- Session management

## Authorization Levels:

### Roles:

- user: Basic authenticated user
- creator: Content creator *with* plot tools
- superadmin: Full system access
- author: Book author privileges

### Access Control:

- Row-Level Security (RLS) *in* database
- User can only access their own resources
- Role-based feature gating
- API rate limiting per tier

## API Security

### Rate Limiting:

#### Tier-based limits:

- Free: 10 requests/hour
- Basic: 100 requests/hour
- Standard: 500 requests/hour
- Premium: 2000 requests/hour
- Professional: Unlimited

### Input Validation:

- Pydantic schemas for all requests
- File upload validation (type, size)
- Text sanitization for AI prompts
- SQL injection prevention (ORM)
- XSS prevention (sanitized output)

## Data Security

### At Rest:

- Supabase encryption (AES-256)
- Secure credential storage
- No plaintext passwords

### In Transit:


- HTTPS/TLS 1.3 for all connections
- Secure WebSocket (WSS)
- Encrypted API keys

### PII Protection:


- Email masking in logs
- GDPR-compliant data handling
- User data deletion on request


## Deployment Architecture

### Current Deployment (Render)



### Recommended Production Architecture





## Environment Configuration

### Development (.env.development):

```
ENVIRONMENT=development
DEBUG=true
DATABASE_URL=postgresql://localhost:5432/litink_dev
REDIS_URL=redis://localhost:6379
FRONTEND_URL=http://localhost:5173
```

### Production (.env.production):

```
ENVIRONMENT=production
DEBUG=false
DATABASE_URL=<supabase-connection-string>
REDIS_URL=<render-redis-url>
FRONTEND_URL=https://litinkai.com
```

## Docker Compose Setup

See [backend/docker-compose.yml](#):

- API service (FastAPI)
- Celery worker
- Redis

- Flower (monitoring)
  - Mailpit (dev email)
- 

## Scaling Strategy

### Horizontal Scaling

#### API Layer:

```
Current: 1 instance
Target: 3–5 instances behind load balancer
Strategy: Round-robin with health checks
Session: Stateless JWT (no sticky sessions)
```

#### Worker Layer:

```
Current: 1 Celery worker
Target: 5–10 workers per queue
Queues:
  - high_priority: 5 workers (paid tiers)
  - normal_priority: 3 workers
  - low_priority: 2 workers (free tier)
Scaling: Based on queue depth
```

#### Database:

```
Current: Single Supabase instance
Target: Read replicas for analytics
Strategy: Write to primary, read from replicas
Connection Pooling: PgBouncer
```

## Vertical Scaling

### API Instances:

- CPU: 2 cores → 4 cores
- RAM: 4GB → 8GB
- Workers: 4 per instance

### Celery Workers:

- CPU: 4 cores (video processing)
- RAM: 8GB (FFmpeg operations)
- Concurrency: 4–8 tasks per worker

## Caching Strategy

### Redis Cache Layers:

L1: Application cache (5 min TTL)

- User session data
- Subscription info

L2: API response cache (1 hour TTL)

- Book metadata
- Plot overviews
- Character profiles

L3: AI response cache (24 hour TTL)

- Generated scripts (by hash)
- Character descriptions
- Scene descriptions

## CDN Strategy

Static Assets: Cloudflare CDN

- Frontend JS/CSS
- Images
- Fonts

Media Files: Supabase Storage + CDN

- Generated videos
- Audio files
- Character images

## Queue Management

### Priority Queues:

```
celery -A app.tasks.celery_app worker \
-Q high_priority,normal_priority,low_priority \
--concurrency=4
```

### Task Routing:

Task Priority:

- professional/premium → high\_priority
- standard/basic → normal\_priority
- free → low\_priority

Retry Policy:

- Max retries: 3
- Backoff: exponential ( $2^{\text{retry}}$  seconds)
- Fallback: Switch to cheaper model

---

## Documentation Index

### Architecture Documents (This Folder)

1. [README.md](#) - This document (overview)
2. [C4-DIAGRAMS.md](#) - C4 model (Context, Container, Component, Code)
3. [UML-DIAGRAMS.md](#) - UML diagrams (Class, Sequence, Activity)
4. [ERD.md](#) - Entity Relationship Diagram
5. [ADR-INDEX.md](#) - Architecture Decision Records
6. [API-SPECIFICATIONS.md](#) - API documentation
7. [DEPLOYMENT.md](#) - Deployment and infrastructure
8. [TECH-STACK.md](#) - Detailed technology choices

### Existing Architecture Docs (./architecture-docs/)

- [AI\\_PLATFORM\\_ARCHITECTURE.md](#)
- [CELERY\\_ARCHITECTURE.md](#)
- [PLOT\\_GENERATION\\_ARCHITECTURE.md](#)
- [SCRIPT\\_SYNCHRONIZATION\\_ARCHITECTURE.md](#)
- [litinkai\\_technical\\_guide.md](#)

### Implementation Guides

- [OPENROUTER\\_IMPLEMENTATION\\_GUIDE.md](#)
- [EMAIL\\_VERIFICATION\\_SETUP.md](#)
- [CELERY\\_SETUP\\_INSTRUCTIONS.md](#)
- [SUPABASE\\_STORAGE\\_SETUP.md](#)

---

## Quick Start

### Prerequisites

- Python 3.11+
  - Node.js 18+
  - Docker & Docker Compose
  - Redis
  - PostgreSQL (or Supabase account)

### Backend Setup

```
cd backend
python -m venv venv
source venv/bin/activate
pip install -r requirements.txt
cp .env.example .env
# Configure environment variables
docker-compose up -d redis
uvicorn app.main:app --reload
```

## Frontend Setup

```
npm install
cp .env.example .env
# Configure environment variables
npm run dev
```

## Start Workers

```
cd backend
celery -A app.tasks.celery_app worker --loglevel=info
```

---

## Support & Contact

- **Documentation:** See [Documentation Index](#)
  - **Issues:** GitHub Issues
  - **Email:** support@litinkai.com
- 

**Last Updated:** 2025-11-06

**Version:** 2.0

**Status:** Production Ready